

# Akka for Java Developers

# Unit 1. What is Akka?

Akka for Java Developers

## What is Akka?

Definition from <http://akka.io>:

- Akka is a toolkit and runtime for building highly concurrent, distributed, and fault tolerant event-driven applications on JVM.

Akka is written in Scala but has API for Java.

Akka's source is open.

## Toolkit and runtime

Akka gives you both API and small-footprint runtime.

Footprint is less than 10MB.

## Highly concurrent applications

- Modern applications handle multiple requests in parallel. Traditional parallel computing has many issues.
- Akka implements Actor Model for concurrency.
- It abandons shared memory model therefore removes the need for blocking.
- Unfortunately in Java we can abandon shared memory only by convention. There is no language support like in Erlang.

## Fault-tolerant applications

- Modern applications are deployed in complex environments and communicating with many data sources and external systems.
- Too many points of failure. It is hard to make your application stable enough.
- Akka gives a well-proved mechanism of supervision that allows you to build self-healing systems.

**“Let it crash”** is the main idea behind this.

## Event-driven applications

Akka gives you a simple asynchronous model of communication which resembles how things and people interact in real world.

## Distributed applications

- Akka effectively implements Location Transparency for each building block.
- Distribution and scaling is very easy and transparent.
- It does not require any major changes in code for Akka application to become distributed.

## Applications for JVM

Akka was created in Scala but has Java API.  
This makes it useable on any JVM language.

# Unit 2. Foundation


Akka for Java Developers

## Problem of Concurrency


Main problem:

**Competition over the mutable state**


This is your mutable shared state




Multiple tasks would like to access it



They do this in parallel




Problem 1






To preserve consistency these updates are applied sequentially



By placing everyone in a queue




First one to get a lock wins



Other tasks are waiting in a queue




Which is not a queue




Problem 2



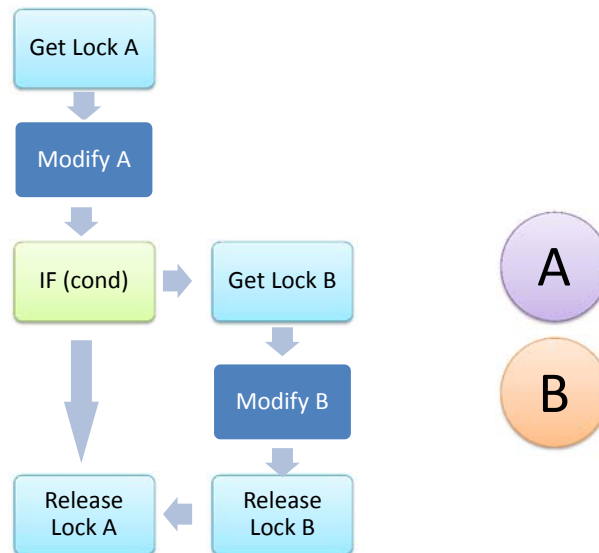
## Locks?



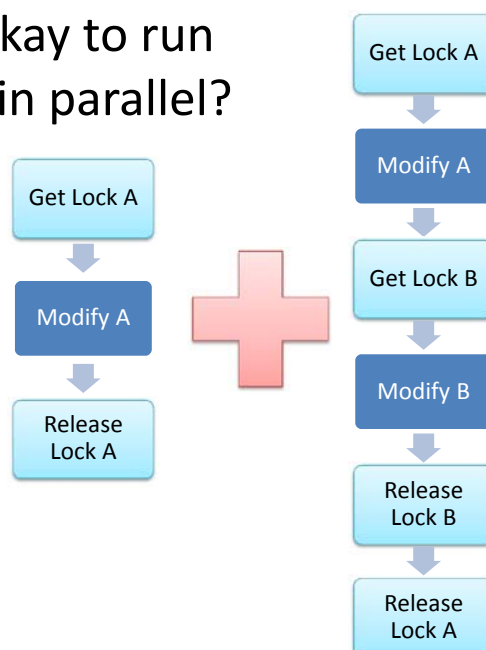
Yes, for very simple cases



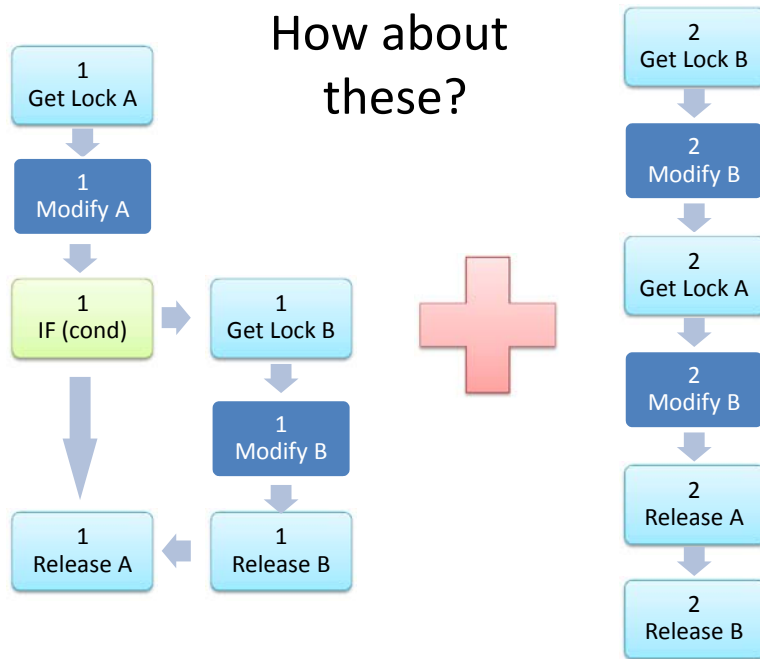
But bad for multiple resources



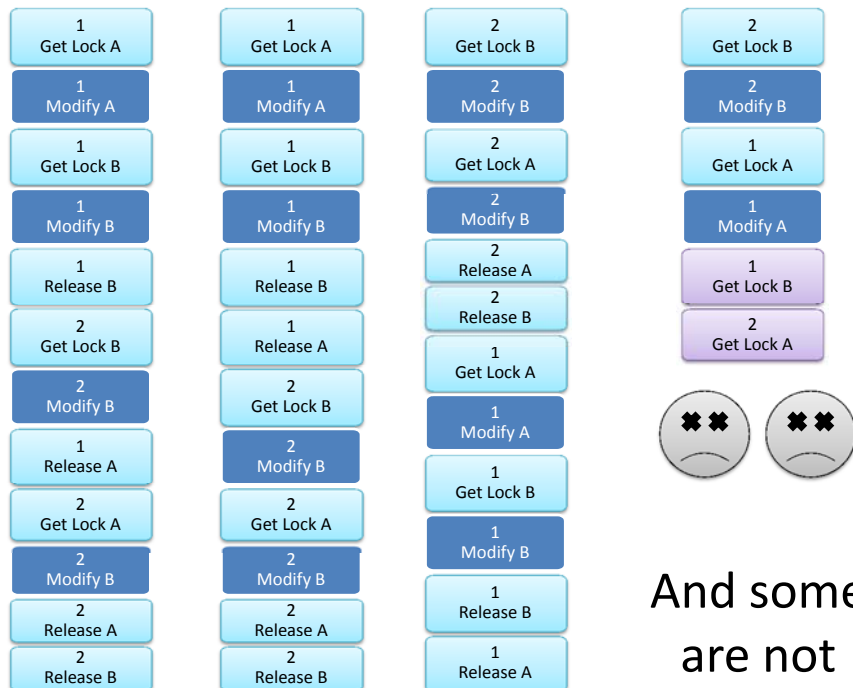
Is it okay to run these in parallel?



How about these?




Some of possible scenarios are okay



And some  
are not

## Problem 3




## Actor Model

- Everything is an actor
- No shared state
- Message passing

**Actor =**  
mailbox + behavior + state



Actors receive messages



Actors are lightweight



# Actors

Actor is the universal primitive  
of concurrent computation.

It collects messages from mailbox  
and reacts to them.

## Local Decisions

It can make local decisions:

- modify private state
- create new actors
- send messages to other actors
- determine how to respond to the next message received.

## Modify private state



## Create new actor



Send messages to other actors



Determine how to respond to the next message received



## Local Effects

- All effects that are produced by actors are local. Actor only can affect on things about which it is aware.
- There is also no simultaneous change in multiple locations. No mutable shared state is permitted.

## Function calling vs Message passing

### **Function calling:**

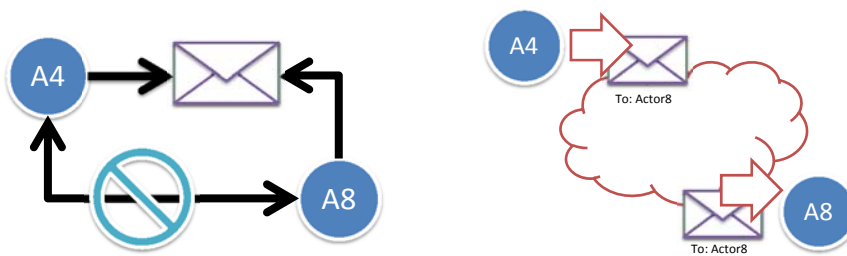
- It requires the reference and the interface.
- This creates unnecessary dependencies and introduces many limits to possible designs.
- It also supports only synchronous communication.

### **Message passing :**

- It requires only the address to which the message is sent.
- There is no need to acquire a reference to the recipient.
- There is no need to interact with recipient directly at all.
- This allows building systems with variable topology.

## Message passing

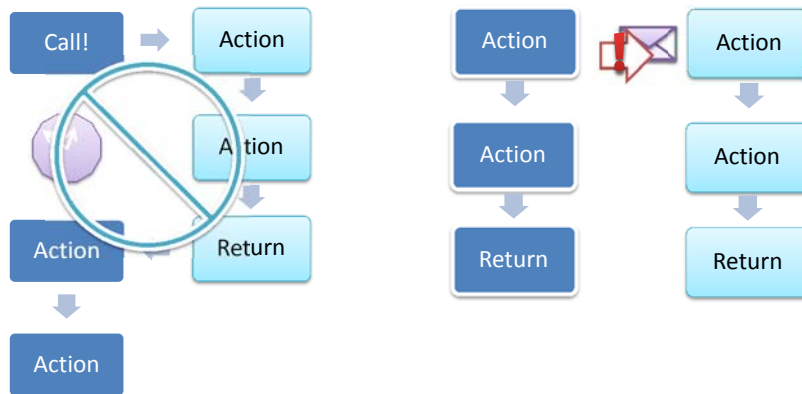
- The only proper way of communication within actor-based applications is through passing messages.
- No mutable messages are permitted.



## Non-blocking behavior

- Locks and synchronized blocks simulate “time freezes”.
- Actor model allows us to embrace the time.
- In general it is not required for actors to use only non-blocking behavior but in heavy-loaded systems blocking causes issues with performance.
- There almost always is a way to solve the problem in a non-blocking style.

## Non-blocking behavior



## Akka Actor Example

```
import akka.actor.UntypedActor;

public class SimpleActor extends UntypedActor {

    @Override
    public void onReceive(Object message)
        throws Exception {
        if("Hello!".equals(message)){
            System.out.println("Oh, hi there!");
        }
    }
}
```

# Akka Actor Application example

```
import akka.actor.ActorRef;
import akka.actor.ActorSystem;
import akka.actor.Props;

public class AkkaExample {
    public static void main(String[] args) throws Exception {

        ActorSystem system = ActorSystem.create("Example");

        Props actorProps = Props.create(SimpleActor.class);
        ActorRef actor = system.actorOf(actorProps);

        actor.tell("Hello!", null);

        Thread.sleep(100);

        system.shutdown();
    }
}
```

## ActorSystem

```
ActorSystem system = ActorSystem.create("Example");
```

- ActorSystem represents the environment in which actors are running. Treat it like a logical application instance.
- Each actor system creates a set of threads and execution contexts on which your application will be executed.
- It is possible to run multiple actor systems but you should know that each one is heavyweight enough so be careful.



# Props

```
system.actorOf( Props.create( SimpleActor.class ) )
```

- Props class is a container which describes how the Actor should be created.
- In general it should contain at least the class of the actor, and may also contain constructor arguments and some other information.

# ActorRef

```
ActorRef actor = system.actorOf(actorProps);  
actor.tell("Hello!", null);
```

- Actor reference is a handle to the actor instance.
- You only can interact with actor through the ActorRef.
- It is **immutable**, **serializable** and **network-aware** so it is safe to pass it to other actors.

# UntypedActor

```
public class SimpleActor extends UntypedActor {  
  
    public void onReceive(Object message) throws Exception {  
        ...  
    }  
}
```

- This is a base class for creating actors in Java.
- You have to implement onReceive method.

## Starting an actor revisited

```
Props actorProps = Props.create(SimpleActor.class);  
ActorRef actor = system.actorOf(actorProps);
```

- Actor can be named.

```
ActorRef actor = system.actorOf(actorProps, "Name");
```

- Name should not start with '\$' sign and must be unique within one level of hierarchy.


## More Tools for Actors

- UntypedActorContext**
- This is an object which gives you more controls for actor's lifecycle and access to various ActorSystem facilities.
  - This object has its own actorOf methods which gives you an opportunity to launch new actors from within your actor.
  - You can get your actor's context object by calling `getContext()`.
- ActorContext**
- Is a more general form of ActorContext.
  - It is the same object but has only the essential Scala API.
  - You can get it by calling `context()`.

## ActorSystem structure

- Actors form hierarchies.
- Your actor is always created as child of another actor.
- When starting actor system creates the root of hierarchy – guardian actor (named “user”).

## Actor Hierarchies



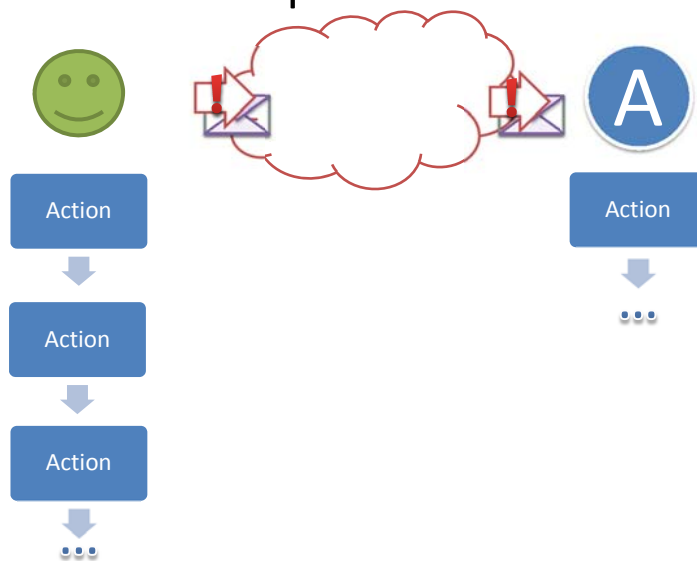
## Actor Hierarchies

- **ActorSystem.actorOf** creates your actors as children of the root actor.
- Your actors can create their own children by calling **ActorContext.actorOf** methods.
- Actor names which you optionally supply to **actorOf** method must be unique only among sibling actors.

# Unit 3. Futures

Akka for Java Developers


It is easy to launch parallel task if you do not depend on its result



## What is a Future?

- Future is a simple way to efficiently perform simultaneous operations.
- It is a container which contents will normally become available sometime after creation.
- Future may contain a result of computation or an exception raised during computation.

## Future Lifecycle



# How Future Works

Future is a collection of 0 or 1 elements.

Future can be created:

- **Empty** → It is pending for a value to appear
- **With value** → future is complete upon creation
- **With exception** → future is failed upon creation

Futures are essentially non-blocking.

## Launching a parallel task with Akka

```
import akka.actor.ActorSystem;
import scala.concurrent.Future;
import akka.dispatch.Futures;
import java.util.concurrent.Callable;

final ActorSystem system = ActorSystem.create();

Future<String> pending = Futures.future(
    new Callable<String>() {
        @Override
        public String call() throws Exception {
            return "Hello from the Future!";
        }
    }, system.dispatcher());
```

## Futures API basics

- `scala.concurrent.Future` is a Scala interface.
  - You can get objects with this interface from many places. From client perspective they all behave the same but they may work differently inside.
  - There is also an “old” interface in `scala.parallel` package. Try not to get confused.
- `akka.dispatch.Futures` is a Akka’s Java-friendly interface for creating futures.

## Creating Future

```
Future<String> pending =  
    Futures.future(callable, dispatcher);
```

```
Future<String> successful =  
    Futures.successful("hello");
```

```
Future<String> failed =  
    Futures.failed(exception);
```



## Consuming Future Values

Future values can be consumed **asynchronously** or **synchronously**.

- Though asynchronous consumption is encouraged, there are cases when it is not possible to use it.
- Synchronous consumption may cause a current thread to block.

## Consuming asynchronously

Asynchronous consumption is similar to consuming events. It requires an event handler.

Event subscription	Result is delivered to
<code>Future&lt;T&gt;.onSuccess</code>	<code>OnSuccess&lt;T&gt;.onSuccess(T value)</code>
<code>Future&lt;T&gt;.onFailure</code>	<code>OnFailure.onFailure(Throwable e)</code>
<code>Future&lt;T&gt;.onComplete</code>	<code>OnComplete&lt;T&gt;.onComplete(Throwable e, T value)</code>

You can attach multiple event handlers to the Future

## Consuming asynchronously

```
pending.onSuccess(new OnSuccess<String>() {  
    @Override  
    public void onSuccess(String result)  
        throws Throwable {  
        System.out.println("Yay! I got " + result);  
    }  
}, system.dispatcher());
```

## Consuming synchronously

- **Await** class helps in waiting for Future results.
- **It blocks.** But it never blocks forever.
- With this model deadlocks are not possible.

## Consuming synchronously

```
import scala.concurrentAwait;
import scala.concurrent.durationDuration;
import java.util.concurrent.TimeUnit;

Duration duration = Duration.create(5, TimeUnit.SECONDS);

String value = Await.result(pending, duration);

//Other ways to create durations
Duration duration2 = Duration.create(5, "seconds");
Duration duration3 = Duration.create("5 seconds");
```

## Programming the Future

- If the result of parallel task must be transformed after being received, use **mapping**.
- If there are many parallel operations but they are all needed to calculate the final result, use **sequence**, **fold** or **reduce**.
- You can chain these operations.

## Mapping the result

```
Future<String> pending = Futures.future(callable, dispatcher);

Future<Integer> mapped = pending.map(

    new Mapper<String, Integer>() {
        @Override
        public Integer apply(String value) {

            return value.length();

        }
    }, system.dispatcher());
```

## Sequence and Reduce Futures

```
List<Future<String>> listOfFutureStrings= ...

Future<Iterable<String>> futureListOfStrings=
    Futures.sequence(
        listOfFutureStrings,
        system.dispatcher());

Future<Integer> futureInteger = Futures.reduce(
    listOfFutureStrings,
    new akka.japi.Function2<Integer, String, Integer>() {
        @Override
        public Integer apply(Integer arg1, String arg2) {
            return arg1 + arg2.length();
        }
    }, system.dispatcher());
```

# Unit 4. Creating Actors

Akka for Java Developers

## Akka Actor Example

```
import akka.actor.UntypedActor;

public class SimpleActor extends UntypedActor {

    @Override
    public void onReceive(Object message)
        throws Exception {

        if("Hello!".equals(message)){
            System.out.println("Oh, hi there!");
        }
    }
}
```

## Starting Actors

Actors are started at some point of hierarchy:

- `ActorSystem.actorOf(Props props)`  
Will start a root actor.
- `ActorContext.actorOf(Props props)`  
Will start a child actor.

## Props tell how actor is created

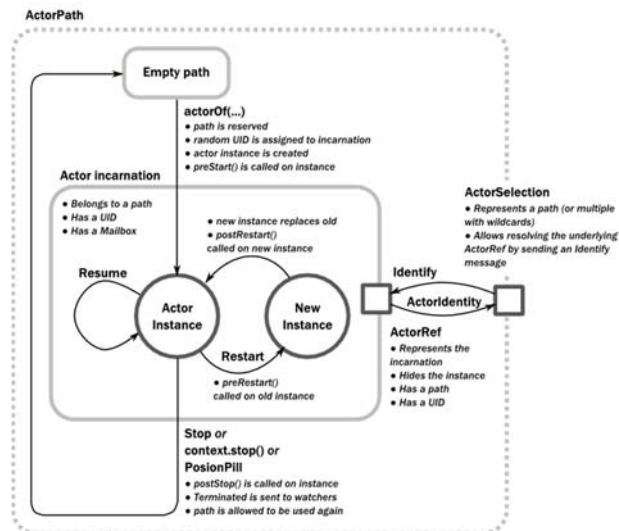
Props are describing how to create an actor:

```
Props props = Props.create(MyUntypedActor.class)

Props props = Props.create(
    MyUntypedActor.class, arg1, arg2)

Props props = Props.create(
    new Creator<MyActor>() {
        @Override public MyActor create() {
            return new MyActor("...");
        }
    });
```

# Actor Lifecycle



## Lifecycle methods

You can override these methods:

Method	Purpose
<code>preStart()</code>	To perform operations after constructor fired and exactly before actor can receive its first message.
<code>postStop()</code>	Is called after the actor stops receiving messages and before the actor is detached from its path.
<code>preRestart()</code>	Is called on failed instance before it becomes replaced.
<code>postRestart()</code>	Is called on new instance in the same moment as <code>preStart()</code> .

## onReceive method

- Messages from actor's mailbox are delivered to **onReceive** method one by one.
- Actor may react on messages.
- It may detect message type or identity by using `instanceof` or `equals()`.
- There are few messages that are processed by actor itself and never reach onReceive.

## Special messages

- Sending **PoisonPill** will result in stopping the recipient actor.
- Upon receiving an **Identify** message actor responds with an **ActorIdentity** message. It can be used to find ActorRefs in system.



## Actor's state

- Actor's **onReceive** method never runs in parallel. It is very safe to change actor's private state without locks or synchronize blocks.
- You can create an initial state in the constructor or in **preStart()** method.
- You must never share the state.

## Communicating to Actors

There is only one way to communicate with an actor: `ActorRef.tell(message, sender)`

**sender** parameter allows you to attach a return address to your message so recipient can send an answer. You may pass **null** value if you are not expecting any reply.

## Communicating to Actors

**Ask** is a convenient pattern for cases when you are expecting a reply from an actor.

```
Future<Object> future = Patterns.ask(actor, msg, timeout);
```

Timeout indicates for how long this Future will wait for a reply before failing. You never wait forever.

## How actor replies

Actor replies by telling to a **ActorRef** from **getSender()**.

This **ActorRef** points to a dead letters mailbox if there is no sender information.

```
public void onReceive(Object message) {  
    if("Hello!".equals(message)){  
        getSender().tell("Oh, hi there!", getSelf());  
    }  
}
```

## Dead Letters Mailbox

**ActorSystem** has a special mailbox to which all undeliverable messages are redirected.

**ActorSystem.deadLetters()** gives you a reference to this mailbox.

## Stopping Actors

Three ways:

- `ActorSystem.stop(ActorRef ref)`
- `ActorContext.stop(ActorRef ref)`
- `ActorRef.tell(PoisonPill.getInstance(), null);`

## Stopping Actors

Stop procedure is performed asynchronously:

- Current message processing completes.
- Actor suspends mailbox.
- All next messages are gone to Dead Letters.
- Actor sends stop commands to its children.
- It awaits for all children to stop.
- It announces its termination to supervisors.
- It finally stops.

## Stop() vs PoisonPill

- `stop()` gives the actor a chance to process current message and then shuts it down
- `PoisonPill` is placed to the end of the mailbox queue. Actor will shut down after it reaches this message.

Anyway, it takes time to shut down an Actor.

# Unit 5. Wiring Actors

Akka for Java Developers

## How actors discover each other

Actor only can communicate with another actor if it knows something about it:

- ActorRef
- Path

## References and Paths

- **Actor Path** represent an address in some ActorSystem which may or may not be inhabited by an actor at some point of time.
- **Actor Reference** points to a specific incarnation of an actor which temporarily occupies some Path.
  - When actor stops the reference becomes invalid and messages will not be delivered to the new incarnation on the same Path.

## Obtaining ActorRef

ActorRef may be obtained:

- Via constructor parameter
  - Pass it through the Props instance
- By receiving in message
  - ActorRefs are immutable and serializable and can be passed freely or transferred over a network
- By creating a child actor
  - You get ActorRef and then can obtain a list of children from `getContext().children()`

## Discover Actor by its Path

ActorSystem and ActorContext can convert Path to ActorSelection.

- ActorSelection allows communicating with an actor behind a Path.
- It does not tied to any actor incarnation.
- The actor corresponding to the selection is looked up when delivering each message.

## Invalid destinations

Messages to invalid destinations such as:

- an ActorRef to a dead Actor
- an ActorSelection to an unoccupied path

are delivered to Dead Letters Mailbox.

# Unit 6. Messages and State

Akka for Java Developers

## Java is Java

- There is no language support for non-shareable state.
- There is no language support for enforcing state immutability
- There is a language support for sharing and breaking every piece of state possible.



# Immutable messages

Design messages to be immutable.

- Use final fields
- Hide set-methods
- Use builders

## Why immutable messages?

- Mutable messages are a form of shared state.
- Without locks there is a chance of receiving message in inconsistent state.
- Also you never know how this message will travel:
  - It may go through the network
  - It may go through variable set of intermediaries

## Immutable state

- If your actor has a state it may need to expose it through sending it in messages
- Sending a state in messages directly is an **act of sharing**.
- If you want to send your state in messages, it **should be immutable**.
- If you cannot design your state immutable, then you **must send copies** of your state.

## State and Parallel tasks

If you are using Futures for processing something in actors, then you must be careful.

- Parallel tasks often are executed, well, in parallel.
- Futures launched during processing one message may complete after the actor moved on to a different messages.

**Watch your captures!**

# Beware!

```
public void onReceive(Object message){
    if (message instanceof SomeMessage) {
        //Do something
        Futures.future(new Callable<Object>() {
            public SomeResult call() {
                //do something in parallel
                getSender().tell(something, getSelf());
                return null;
            }
        }, context().system().dispatcher());
    }
}
```

# Save all that might change

```
public void onReceive(Object message){
    if (message instanceof SomeMessage) {
        //Do something
        final ActorRef sender=getSender();
        Futures.future(new Callable<Object>() {
            public SomeResult call() {
                //do something in parallel
                sender.tell(something, getSelf());
                return null;
            }
        }, context().system().dispatcher());
    }
}
```

# Just don't.

```
int counter=0;

public void onReceive(Object message){
    if (message instanceof SomeMessage) {
        //Do something
        Futures.future(new Callable<Object>() {
            public SomeResult call() {
                //do something in parallel
                counter++;
                return null;
            }
        }, context().system().dispatcher());
    }
}
```

## Possible solution

```
int counter=0;

public void onReceive(Object message) throws Exception {
    if (message instanceof SomeMessage) {
        Futures.future(new Callable<Object>() {
            public SomeResult call() {
                //do something in parallel
                getSelf().tell(new IncrementCounter(),null);
                return null;
            }
        }, context().system().dispatcher());
    }
    if(message instanceof IncrementCounter){
        counter++;
    }
}
```